


Deep Learning with Python: MNIST Example


Deep Learning for MNIST Data

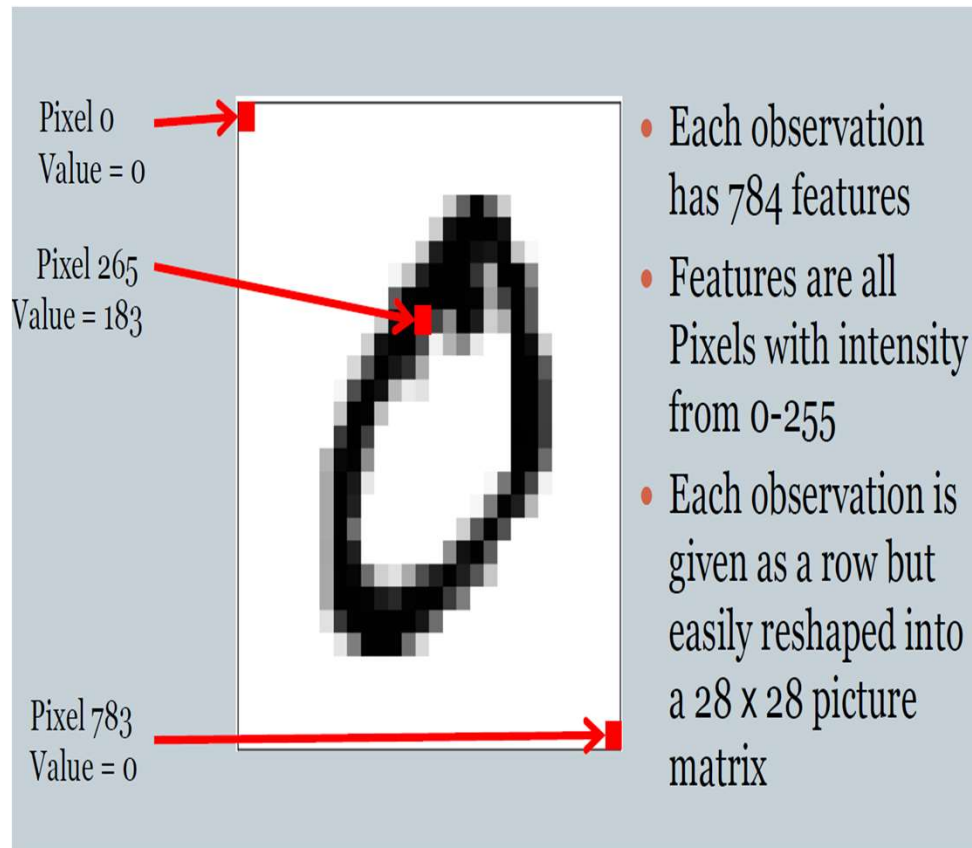


MNIST Data

- The MNIST database (Modified National Institute of Standards and Technology database) of handwritten digits consists of a training set of 60,000 examples, and a test set of 10,000 examples.
- It is a subset of a larger set available from NIST. Additionally, the black and white images from NIST were size-normalized and centered to fit into a 28x28 pixel bounding box and anti-aliased, which introduced grayscale levels.



Each observation represents an image of a handwritten digit 0-9:






- Each image is 28 pixels in height and 28 pixels in width, for a total of 784 pixels in total.
- Each pixel has a single pixel-value associated with it, indicating the lightness or darkness of that pixel, with higher numbers meaning darker. This pixel-value is an integer between 0 and 255, inclusive.

How does a computer read an image?



0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0
0	0	0	0	1	12	0	11	39	137	37	0	152	147	84	0	0	0	0	0
0	0	1	0	0	0	41	160	250	255	235	162	255	238	206	11	13	0	0	0
0	0	0	16	9	9	150	251	45	21	184	199	154	255	233	40	0	0	0	0
10	0	0	0	0	0	145	146	3	10	0	11	124	253	255	187	0	0	0	0
0	0	3	0	4	15	236	216	0	0	38	109	247	240	169	0	11	0	0	0
1	0	2	0	0	0	253	253	23	62	224	241	255	164	0	5	0	0	0	0
6	0	0	4	0	3	252	250	228	255	255	234	112	28	0	2	17	0	0	0
0	2	1	4	0	21	255	253	251	255	172	31	8	0	1	0	0	0	0	0
0	0	4	0	163	225	251	255	229	120	0	0	0	0	0	11	0	0	0	0
0	0	21	162	235	255	254	255	126	6	0	10	14	6	0	0	9	0	0	0
3	79	242	255	141	66	255	245	189	7	8	0	0	5	0	0	0	0	0	0
26	221	237	98	0	67	251	255	144	0	8	0	0	7	0	0	11	0	0	0
125	255	141	0	87	244	255	208	3	0	0	13	0	1	0	1	0	0	0	0
145	248	228	116	235	255	141	34	0	11	0	1	0	0	0	1	3	0	0	0
85	237	253	245	255	210	21	1	0	1	0	0	6	2	4	0	0	0	0	0
6	23	112	157	114	32	0	0	0	0	2	0	8	0	7	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0

Handwriting Digit Recognition



In deep learning, the function f is represented by neural network



Pixel data of an image

- For the 28×28 pixel images we've been using, this means our network has 784 ($=28 \times 28$) input neurons.
- Each images is then encoded in a $m=785-1=784$ feature vector, representing the intensity of the pixel. Each element of the vector contains an integer encoding the intensity of a pixel. In other words the original image is decomposed in a 28×28 pixel grid, which is reshaped to be represented as a $28 \times 28 = 784$ column vector.
- We will train the network's weights and biases so that the network's output would - we hope! - correctly identify the input image: '0', '1', '2', ..., '8', or '9'.



Training data


- The training data set has 785 columns. The first column, called "label", is the digit that was drawn by the user. The rest of the columns contain the pixel-values of the associated image.
- Each pixel column in the training set has a name like `pixelx`, where x is an integer between 0 and 783, inclusive. To locate this pixel on the image, suppose that we have decomposed x as $x = i * 28 + j$, where i and j are integers between 0 and 27, inclusive. Then `pixelx` is located on row i and column j of a 28×28 matrix, (indexing by zero).
- For example, `pixel31` indicates the pixel that is in the fourth column from the left, and the second row from the top, as in the ascii-diagram below.



What's Softmax Function?

$$f_i(x) = \frac{\exp(x_i)}{\sum_j \exp(x_j)}$$

Example: Softmax layer as the output layer





Example :

$$2.33 \rightarrow P(\text{Class 1}) = \frac{\exp(2.33)}{\exp(2.33) + \exp(-1.46) + \exp(0.56)} = 0.83827314$$

$$-1.46 \rightarrow P(\text{Class 2}) = \frac{\exp(-1.46)}{\exp(2.33) + \exp(-1.46) + \exp(0.56)} = 0.01894129$$

$$0.56 \rightarrow P(\text{Class 3}) = \frac{\exp(0.56)}{\exp(2.33) + \exp(-1.46) + \exp(0.56)} = 0.14278557$$

- 
- 
- Softmax function is used to convert the numerical output to values in the range $[0, 1]$
 - The output of softmax function can be seen as probability distribution given the output sums up to 1
 - Softmax function is used in multiclass classification methods such as neural networks, multinomial logistic regression, multiclass LDA, Naive Bayes classifier.
 - Softmax function is used to output action probabilities in case of reinforcement learning
 - Softmax function is used as an activation function in the last / final layer of neural network algorithm.

Cross Entropy Loss

This is the most common setting for classification problems. Cross-entropy loss increases as the predicted probability diverges from the actual label.

Mathematical formulation :-

$$CrossEntropyLoss = -(y_i \log(\hat{y}_i) + (1 - y_i) \log(1 - \hat{y}_i))$$



Cross entropy loss


For multiclass classification, we calculate a separate loss for each class label per observation and sum the result.



Accuracy

- For example, if y_{true} is $[1, 2, 3, 4]$ and y_{pred} is $[0, 2, 3, 4]$ then the accuracy is $3/4$ or $.75$. If the weights were specified as $[1, 1, 0, 0]$ then the accuracy would be $1/2$ or $.5$.

- 
- 
- one epoch = one forward pass and one backward pass of all the training examples
 - batch size = the number of training examples in one forward/backward pass. The higher the batch size, the more memory space you'll need.
 - number of iterations = number of passes, each pass using [batch size] number of examples. To be clear, one pass = one forward pass + one backward pass (we do not count the forward pass and backward pass as two different passes).
 - Example: if you have 1000 training examples, and your batch size is 500, then it will take 2 iterations to complete 1 epoch.



verbose: Integer. 0, 1, or 2. Verbosity mode. 0 = silent, 1 = progress bar, 2 = one line per epoch.

By setting verbose 0, 1 or 2 you just say how do you want to 'see' the training progress for each epoch.

`verbose=0` will show you nothing (silent)

`verbose=1` will show you an animated progress bar like this:

[=====]


`verbose=2` will just mention the number of epoch like this:

Epoch 1/10



Avoiding Overfitting

- With thousands of weights to tune, overfitting is a problem
Early stopping (when performance starts dropping)
- Regularization terms added to cost function during training
- Dropout -ignore say 50% of all neurons randomly at each training step
- Works surprisingly well! Forces your model to spread out its learning



```
model = Sequential()  
model.add(Dense(512, activation='relu', input_shape=(784,)))  
model.add(Dropout(0.2))  
model.add(Dense(512, activation='relu'))  
model.add(Dropout(0.2))  
model.add(Dense(10, activation='softmax'))
```


One hot encoding (or dummy variables)

```
+---+-----+
| id| country   |
+---+-----+
| 0| russia    |
| 1| germany   |
| 2| australia  |
| 3| korea     |
| 4| germany   |
+---+-----+
```

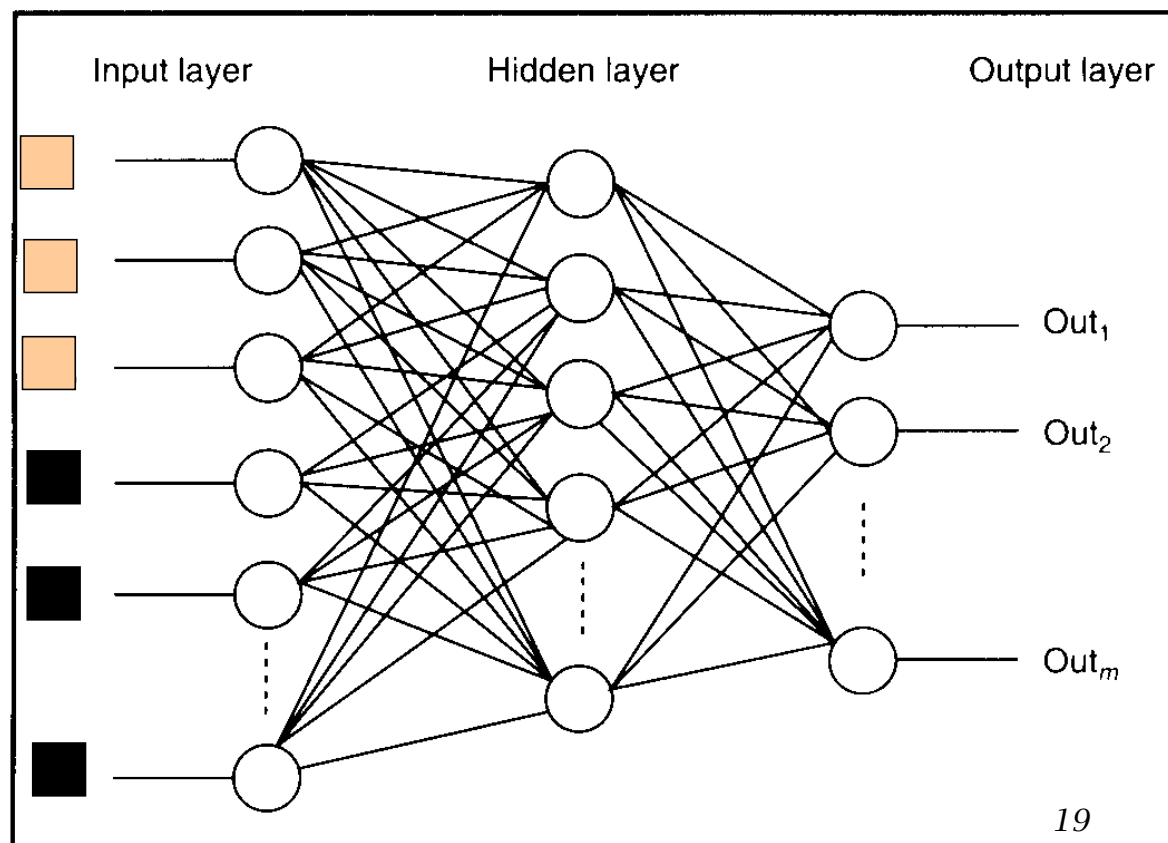
```
+---+-----+-----+-----+-----+
| id| country=russia| country=germany| country=australia| country=korea|
+---+-----+-----+-----+-----+
| 0|          1|          0|          0|          0|
| 1|          0|          1|          0|          0|
| 2|          0|          0|          1|          0|
| 3|          0|          0|          0|          1|
| 4|          0|          1|          0|          0|
+---+-----+-----+-----+-----+
```



Figure 1.2: Examples of handwritten digits from postal envelopes.



Feature detectors








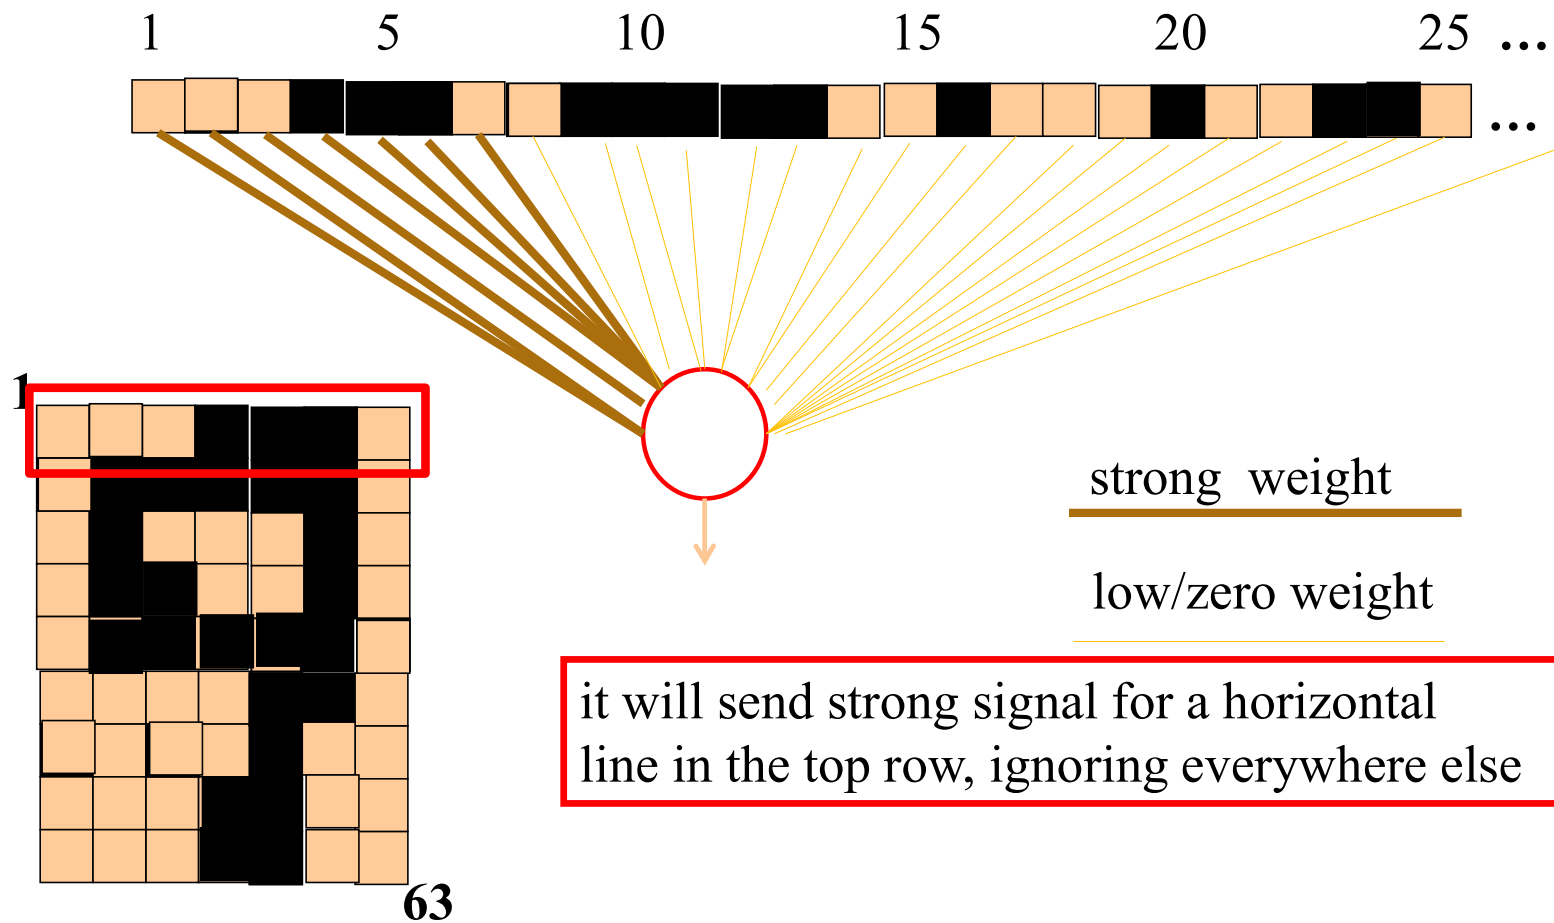
Figure 1.2: Examples of handwritten digits from U.S. postal envelopes.




what is this unit doing?



What does this unit detect?



What does this unit detect?






Figure 1.2: *Examples of handwritten digits from U.S. postal envelopes.*

What features might you expect a good NN to learn, when trained with data like this?

1




Figure 1.2: *Examples of handwritten digits from U.S. postal envelopes.*

1




Figure 1.2: *Examples of handwritten digits from U.S. postal envelopes.*

1




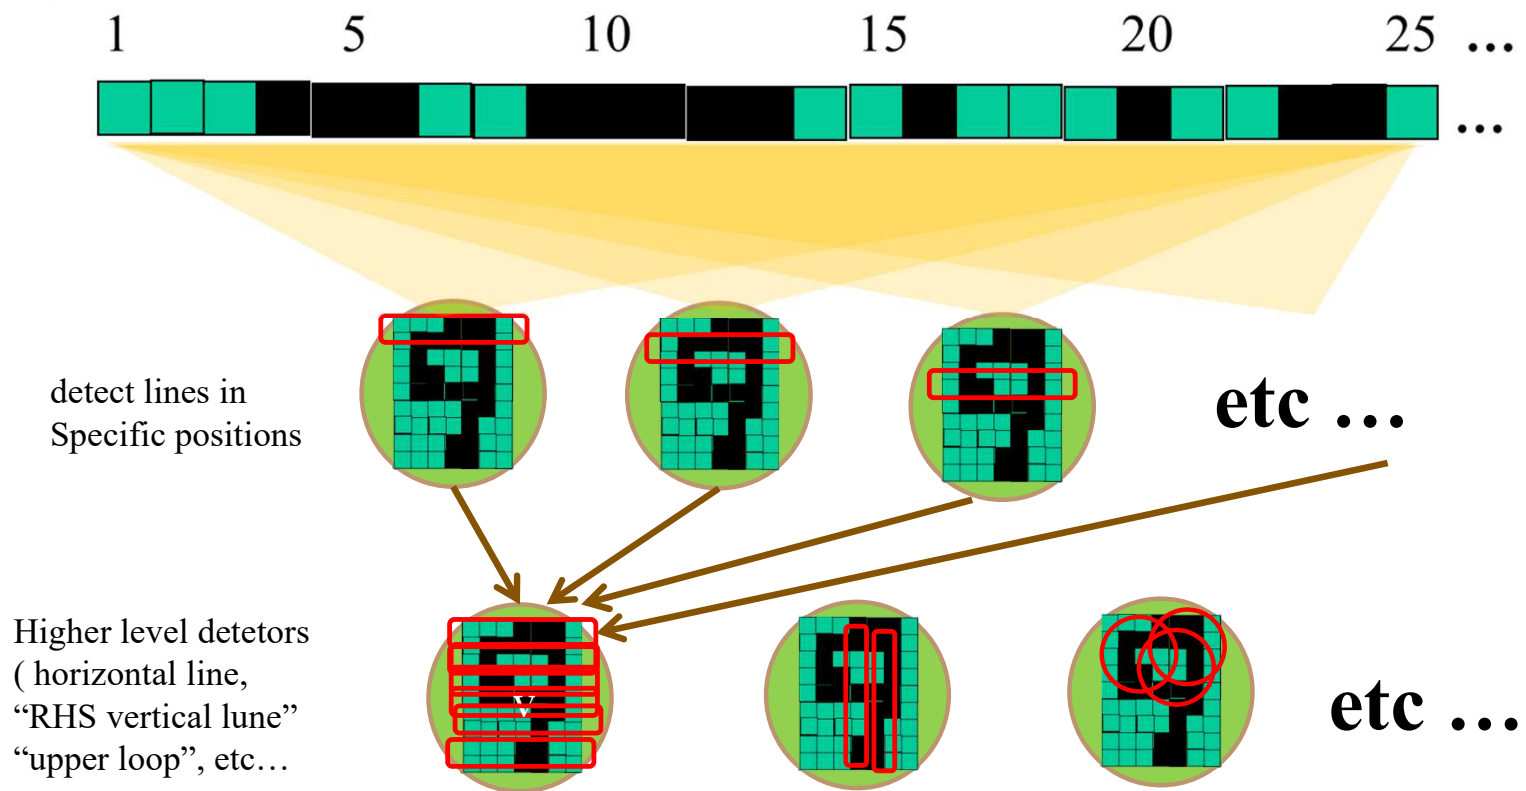
Figure 1.2: *Examples of handwritten digits from U.S. postal envelopes.*

1




But what about position invariance ???
our example unit detectors were tied to
specific parts of the image


successive layers can learn higher-level features ...



successive layers can learn higher-level features ...





So: multiple layers make sense



So: multiple layers make sense

Your brain works that way



So: multiple layers make sense

Many-layer neural network architectures should be capable of learning the true underlying features and 'feature logic', and therefore generalise very well ...

